

Unsupervised Learning - Clustering

Mark Asch - IMU/VLP/CSU

2023

Program

1. Data Analysis

- (a) Introduction: the 4 identifiers of “big data” and “data science”
- (b) Supervised learning methods: regression—advanced, k-NN, linear classification methods, SVM, NN, decision trees.
- (c) **Unsupervised learning methods**: principal component analysis, **k-means, clustering**.

Recall: Introduction

- **Supervised** learning:
 - ⇒ we have p characteristics X_1, X_2, \dots, X_p measured on n observations
 - ⇒ one response Y measured on the same observations
 - ⇒ objective: predict Y using X_1, X_2, \dots, X_p
 - ⇒ methods: regression and classification
- **Unsupervised** learning:
 - ⇒ we **only** have p characteristics X_1, X_2, \dots, X_p measured on n observations
 - ⇒ we want to make predictions, but we do not have an associated response variable...
 - ⇒ objective: **discover interesting effects** with respect to the observations of X_1, X_2, \dots, X_p
 - Can we **visualize** or represent the data in a more **informative** way?
 - Can we **discover sub-groups or clusters** among the variables or the observations?

Recall: Warnings!

- Unsupervised learning is more difficult, since it is subjective.
- ✗ Unsupervised learning is often part of an initial phase of **exploratory data analysis** (EDA), where we compute elementary statistics and plot basic histograms and boxplots---see Introductory lectures.
- ✗ There are no universal **cross-validation** methods for unsupervised learning.
- ✗ There is **no response variable** that can be used to test our models.
- ✓ However, there are a large number of application domains where unsupervised learning is widely used.

Introduction to Clustering

- Clustering is an ensemble of methods for finding subgroups, or clusters, in a database.
- The subgroups are defined according to two principles:
 - ⇒ The observations within each group are **similar**.
 - ⇒ The observations in different groups are **different**.
- For example, suppose we have n observations, each with p properties.
 - ⇒ The properties correspond to **measurements** made for each sample among the n .
 - ⇒ The clustering will enable us to distinguish between **subtypes** within the measurements
 - ⇒ The problem is **unsupervised** because we are attempting to discover the unknown structure, based exclusively on the measured data.

Clustering Approaches

There are a large number of clustering methods. The two best known are:

1. k -means clustering.
2. Hierarchical, or tree clustering.

The major difference is that:

- in k -means we try to divide the observations into a specified number of clusters,
- whereas in the hierarchical approach we do not know in advance the number of clusters, and we end up with a tree-like structure known as a *dendrogram*. This structure enables us to visualize all the possible clusterings, from 1 to n .
- **Note:**

- ⇒ both PCA and clustering seek to simplify the data by means of a small number of parameters, but their **mechanisms** are different:
- ⇒ PCA tries to find a **low dimensional representation** of the observations that explains a large proportion of the variance.
- ⇒ Clustering tries to find **homogeneous subgroups** among the observations.

K-MEANS CLUSTERING

k-means Clustering

- This is a simple and elegant method for dividing a dataset into k distinct clusters with no overlaps.
 - To do the clustering, one must first specify the number of desired clusters, K .
 - The algorithm will then assign each observation to at most one of the K clusters.
- ⇒ Since we do not know the value of k , the algorithm will be based on an optimization problem.




Figure 1: Application of k-means to simulated 2-D data.

Theory of k-means

- Let C_1, \dots, C_K be the sets containing the indices of the observations in each of the K clusters.
- These sets satisfy two properties:
 1. $C_1 \cup C_2 \cup \dots \cup C_K = \{1, \dots, n\}$, so each observation belongs to at least one of the K clusters.
 2. $C_k \cap C_{k'} = \emptyset$ for all $k \neq k'$, meaning that the clusters do not overlap and no observation can belong to more than one cluster.
- Then a “good” clustering is defined as one for which the intra-cluster variation, $W(C_k)$, is the smallest possible. That is, we want to solve the optimization problem

$$\min_{C_1, \dots, C_K} \left\{ \sum_{k=1}^K W(C_k) \right\},$$

with a suitable definition of the variation $W(C_k)$.

- There are several possibilities for this, but we usually just use the **squared Euclidean distance**,

$$W(C_k) = \frac{1}{|C_k|} \sum_{i, i' \in C_k} \sum_{j=1}^p (x_{ij} - x_{i'j})^2,$$

where $|C_k|$ is the number of observations in the k -th cluster.

- Substituting this expression, we obtain the k -means **optimization** problem,

$$\min_{C_1, \dots, C_K} \left\{ \sum_{k=1}^K \frac{1}{|C_k|} \sum_{i, i' \in C_k} \sum_{j=1}^p (x_{ij} - x_{i'j})^2 \right\}. \quad (1)$$

Algorithm for k-means

- Solving this optimization problem is not trivial.
 - ⇒ How do we find a division of the observations into K clusters that guarantees the minimization of the above cost function?
- It can be proved that a very simple algorithm can find an acceptable solution.
 - ⇒ (1) Assign a number between 1 and K at random to each observation--this produces the initial clustering.
 - ⇒ (2) Iterate the following until the cluster assignments do not change anymore:
 - (2a) For each cluster, compute the centroid equal to the vector of the p characteristic means for all observations in cluster p .
 - (2b) Assign each observation to the cluster whose centroid is closest as defined by the chosen distance function.

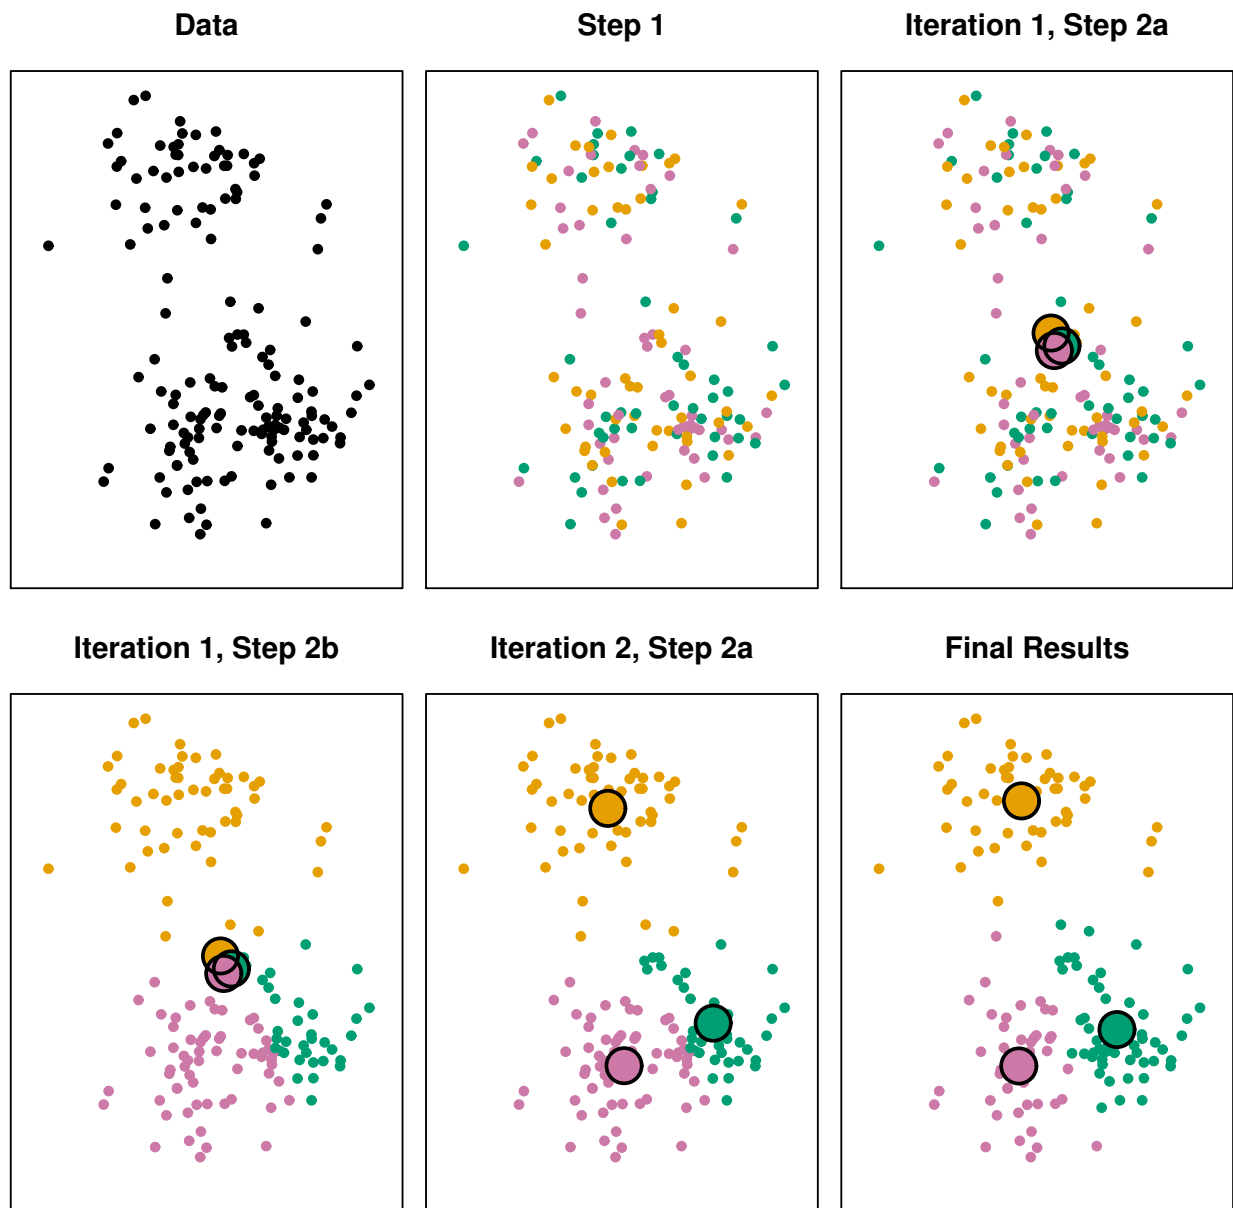


Figure 2: Progress of the k-means algorithm

✓ Since this algorithm finds a **local optimum**, the results

will depend on the initial, random assignment of labels.

- ✓ It is thus vital to relaunch the algorithm a large number of times, starting from different random initial configurations.
- ✓ Finally, we choose the best solution that produces the smallest value of the cost function.



Figure 3: Six runs with $K = 3$

Practical Considerations for Clustering

- Clustering is a very useful tool for data analysis.
- Before starting, we must make two decisions whose potential impacts are fundamental:
 1. Should we normalize (center and reduce) the attributes?
 2. How many clusters should we seek?
- In practice we will have to try several options, since there is no single answer to these two questions.
- The problem of validating the computed clusters is a subject of ongoing research.
- Finally, to interpret the results, one can do the following:
 - ⇒ Tune the parameters—see Examples below.
 - ⇒ Check the robustness by recalculating over subsets of the data—a form of cross-validation.
 - ⇒ Be very cautious when reporting the results—they do not represent the absolute truth!

Example

- The function `kmeans()` performs k-means clustering in R
- Consider a simple, simulated example: there are 2 clusters in the data, with a shift in their mean values between the first 25 and the last 25 observations

```
set.seed(2)
x=matrix(rnorm(50*2), ncol=2)
x[1:25,1]=x[1:25,1]+3
x[1:25,2]=x[1:25,2]-4
head(x)
##           [,1]      [,2]
## [1,]  2.103085 -4.838287
## [2,]  3.184849 -1.933699
## [3,]  4.587845 -4.562247
## [4,]  1.869624 -2.724284
## [5,]  2.919748 -5.047573
## [6,]  3.132420 -5.965878
```

- First, perform k-means with $K = 2$:


```
km.out=kmeans(x,2,nstart=20)
```

- The affectations of the observations are in the variable cluster:

```
km.out$cluster
##  [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
      2 2 2 2 2 2 2 2 2 2 1 1 1 1 1
      1 1 1 1 1
## [36] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

- We observe that k-means has perfectly separated the observations into two clusters, without us having supplied any group information to `kmeans()`...
- Plot the observations, colored by their cluster affectation :

```
plot(x, col=(km.out$cluster+1), main="k-Means Clustering
      Results with K=2", xlab="", ylab="", pch=20, cex=2)
```




- Here, the observations are easy to plot, since we are in 2 dimensions only... In the case of more than 2 variables, we could perform a PCA and then plot the first 2 principal components.
- In this simulated example we knew the number of clusters. But in general this is definitely not the case. So we could have started by trying $K = 3$:

```
set.seed(4)
km.out=kmeans(x,3,nstart=20)
km.out
```

```

plot(x, col=(km.out$cluster+1), main="K-Means Clustering
Results with K=3", xlab="", ylab="", pch=20, cex=2)
## K-means clustering with 3 clusters of sizes 10, 23, 17
##
## Cluster means:
##           [,1]      [,2]
## 1  2.3001545 -2.69622023
## 2 -0.3820397 -0.08740753
## 3  3.7789567 -4.56200798
##
## Clustering vector:
##  [1] 3 1 3 1 3 3 3 1 3 1 3 1 3 1 3 1 3 3 3 3 3 1 3 3 3
##      2 2 2 2 2 2 2 2 2 2 2
## [36] 2 2 2 2 2 2 2 2 1 2 1 2 2 2 2
##
## Within cluster sum of squares by cluster:
## [1] 19.56137 52.67700 25.74089
## (between_SS / total_SS =  79.3 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"
## [5] "tot.withinss" "betweenss"    "size"         "iter"
## [9] "ifault"

```



- Here, the algorithm has divided one of the original 2 clusters...
- Note that to execute `kmeans()` with different initial affectations, we use the argument `nstart`. here we compare `nstart=1` with `nstart=20` by extracting the intracluster sum of squares score.

```
set.seed(3)
km.out=kmeans(x,3,nstart=1)
km.out$tot.withinss
## [1] 104.3319
```

```
km.out=kmeans(x,3,nstart=20)
km.out$tot.withinss
## [1] 97.97927
```

- **Conclusions:**

- ⇒ Restarting has effectively reduced the value of the intracluster sum of squares.
 - ⇒ We have attained a better optimum (minimum).
- It is strongly recommended to execute many times, using a high value of `nstart`, such as 20 or 50, otherwise undesirable local minima will be found...
 - We also use a random seed initialization, by `set.seed()`, to ensure reproducibility of the results.

HIERARCHICAL CLUSTERING

Hierarchical Clustering

- No need to choose the number of clusters **beforehand**.
- Everything is done “ **bottom-up**,” starting from the leaves and moving up to the trunk, or main branch
- This generates a tree-like diagram, known as a **dendrogram**.
- It is a **deterministic** method, but can provide clues for further inference in statistical models.

Interpretation

- Suppose we have synthetic data, in 2D, having 3 distinct classes, but that are observed WITHOUT any knowledge of their classes.




Figure 4: 45 observations in 3 classes




Figure 5: Dendrogram obtained by hierarchical clustering with full linkage.

- Each leaf represents one of the 45 observations (figure on left).
 - Moving up, the leaves merge into branches, corresponding to observations that have a similarity---measured in a given metric---between them.
- ⇒ The earlier the merging takes place, the more similar

are the observations.

⇒ Thus, observations that merge later---towards the top---are more and more different.

- Conclusion: the merging **height**, on the vertical axis, indicates the difference between two observations and it is the height that will determine the number of different clusters.
- Identification of clusters:
 - ⇒ make a horizontal cut across the dendrogram
 - ⇒ the distinct sets below can be interpreted as distinct clusters
 - ⇒ further cuts can be made, going downwards, until each observation is in its own cluster... (this is not very useful!)
- A single dendrogram can provide any number of clusters, from 1 to N , the number of observations.
 - ⇒ Choosing the “good” number of clusters, or cut height, is not an obvious task, and has to be performed visually, based on experience, or familiarity with the data

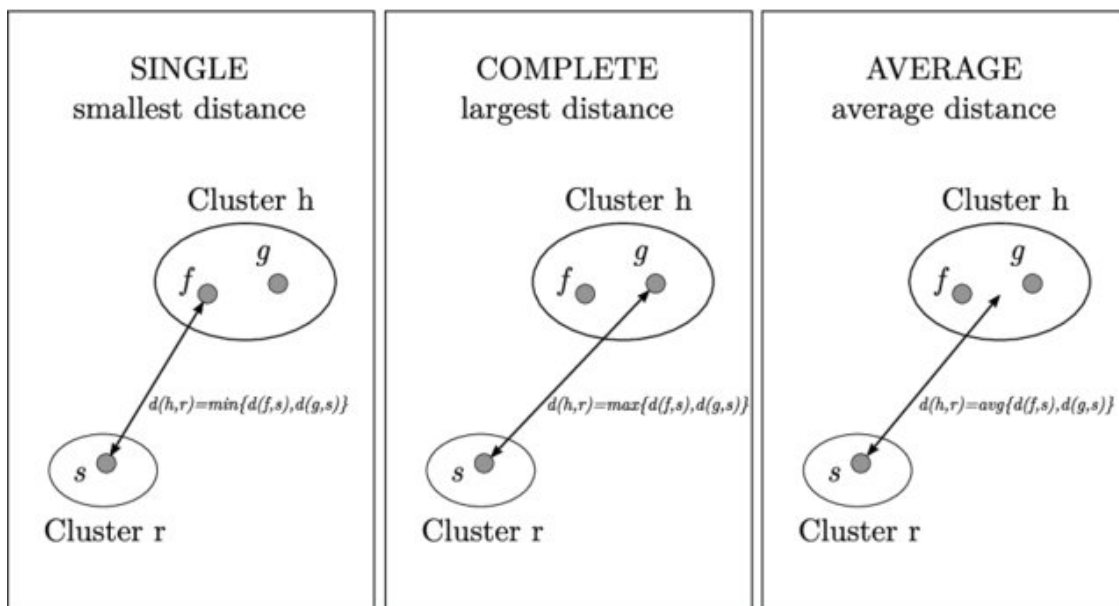
- ⇒ Hopefully, however, there is a clear gap in the link lengths of the dendrogram that will separate the **inherent** clusters from the unnatural ones. This will depend on the context and the data.
- The dendrogram does not work better than K-means, unless the data are truly hierarchical... (without other divisions, eg. male-female)

Algorithm for hierarchical clustering

1. For all n observations, calculate the $n(n - 1)/2$ dissimilarities pairwise. Each observation is a cluster.
2. For $i = n, n - 1, \dots, 2$,
 - (a) calculate all the inter-cluster dissimilarities and identify the pair of two clusters that are the least dissimilar
 - (b) merge these two clusters where the dissimilarity indicates the height in the dendrogram at which the merge takes place
 - (c) calculate the new inter-cluster dissimilarities among the remaining $i - 1$ clusters

Dissimilarity between 2 groups

- Need to generalize the notion of dissimilarity between a pair of observations, to the dissimilarity between a pair of groups (inter-cluster).
- We use the notion of **linkage** that defines the dissimilarity between 2 groups of observations.



- Four types of linkage :

- ⇒ complete
- ⇒ average
- ⇒ single
- ⇒ centroid

- We usually choose **complete** or **average**, since they produce more balanced dendrograms.
- Centroid is often used in genomics.

Further indications

- For the dissimilarity between 2 observations, we use:
 - ⇒ Euclidean distance (most used)
 - ⇒ Correlation distance
- The choice of this distance is very important, and is context-dependent.
- Scaling: as for k-means, we must pay attention to the context.
- It is strongly recommended to test several choices of:
 - ⇒ distance
 - ⇒ linkage
 - ⇒ cut height

Example

- The function `hclust()` performs hierarchical clustering in R.
- Take the following simulated data :

```
set.seed (222)
x=matrix (rnorm (50*2) , ncol =2)
x[1:25 ,1] = x[1:25 ,1] + 3
x[1:25 ,2] = x[1:25 ,2] - 4
head(x)
##           [,1]      [,2]
## [1,]  4.487757 -3.405585
## [2,]  2.998108 -3.479132
## [3,]  4.381021 -4.952044
## [4,]  2.619786 -5.227685
## [5,]  3.184136 -4.202407
## [6,]  2.753104 -2.940880
```

- Use the Euclidean distance to compute the clusterings with complete, average and single linkages.

```

hc.complete = hclust(dist(x), method="complete")
hc.average   = hclust(dist(x), method="average")
hc.single    = hclust(dist(x), method="single")


```

- Plot the dendrograms.

```

par(mfrow=c(1,3))
plot(hc.complete, main="Complete Linkage", xlab="", sub="", cex=.9)
plot(hc.average,  main="Average Linkage", xlab="", sub="", cex=.9)
plot(hc.single,   main="Single Linkage", xlab="", sub="", cex=.9)

```




- Display the cluster labels for each observation associated with a given cut

```
cutree(hc.complete, 2)
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2
[30] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
cutree(hc.average, 2)
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2
[30] 2 2 2 1 2 2 2 2 2 2 2 2 2 2 2 1 2 1 2 2 2 2
cutree(hc.single, 2)
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1 1 1
[30] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

- The complete and average linkages give good results, but the simple is not accurate.
- For scaling, we can use the function `scale()`

```
xsc=scale(x)
plot(hclust(dist(xsc), method="complete"),
     main="Hierarchical Clustering with Scaled Features")
```

Hierarchical Clustering with Scaled Features



References

1. M. DeGroot, M. Schervish, *Probability and Statistics*, Addison Wesley, 2002.
2. Spiegel, Murray and Larry Stephens, *Schaum's Outline of Statistics*, 6th edition, McGraw Hill. 2017.
3. G. James, D. Witten, T. Hastie, R. Tibshirani. *An Introduction to Statistical Learning with Applications in R*. Springer. 2013.
4. T. Hastie, R. Tibshirani, J. Friedman. *The Elements of Statistical Learning*. Springer. 2009.
5. Rachel Schutt and Cathy O'Neil. *Doing Data Science*. O'Reilly. 2014.